


Decentralized Timeline Service

Large Scale Distributed Systems

A84462 | Alexandre Miranda
A84961 | Alexandre Ferreira
A85227 | João Azevedo
A85315 | Miguel Cardoso
A85729 | Paulo Araújo




Universidade do Minho
Escola de Engenharia




System model

- Each node has **an unique ID**
- Published messages **must reach every active subscriber**
- The timeline should describe a **total order of posts**
- Only **gracious faults** are considered
- Decentralization **only relative** to the timeline posts
- All information stored is **ephemeral**




How can we design this system?

Network design




Storage design



Our approach

- **Centralized** and **Decentralized** components
- **Network management** (central servers):
 1. Authentication services
 2. Subscription services
 3. Connection management
- **Network peers** (nodes):
 1. Pub-Sub pattern
 2. Timeline management
 - a. posts ordering and recovery services




Centralized portion

It's used for:

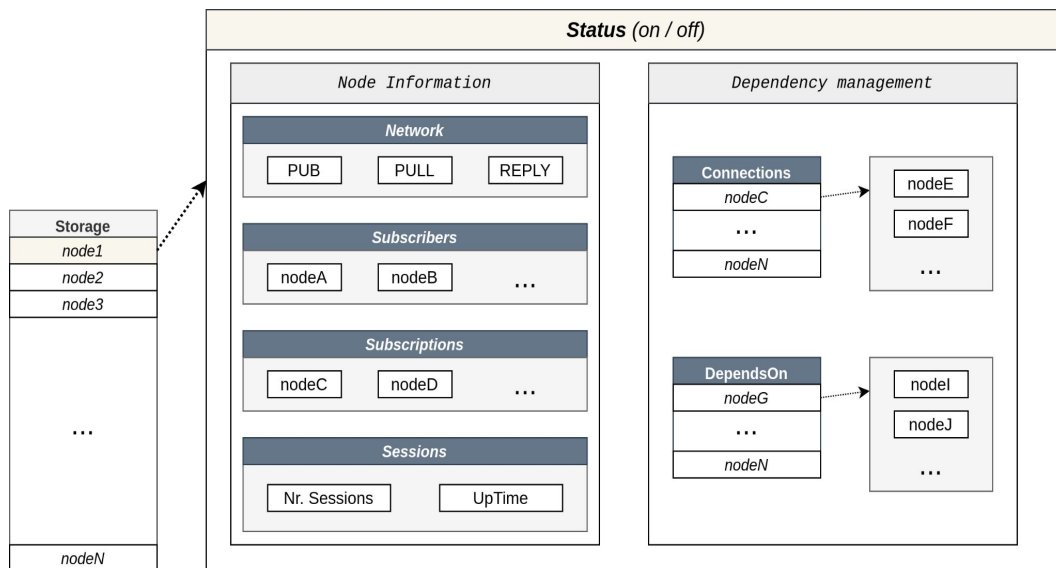
- a. Node maintenance
- b. Network management
- c. Check node status
- d. Auth. and Sub. services

"Entry door to the network"



Node Storage Redis Key-Value Store

Use cases Event-Driven server



1. Register

- a. registers new session

2. Login

- a. new session, manages subscriptions/subscribers

(we will talk about this one later on)

3. Logout

- a. updates session stats
- b. manages dependencies

4. Subscription

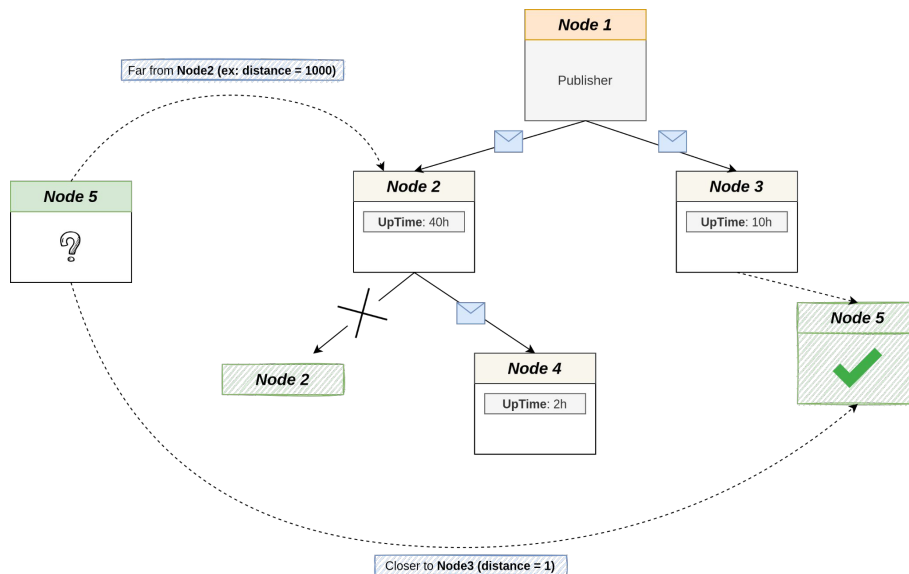
- a. node election

Node election

1. Important for a **reliable subscription service**
2. Candidate with a better score wins
3. How we compute the score?

$uptime = \text{upTime}(\text{candidate})$
 $distance = \text{diff}(\text{candidate.port}, \text{subscriber.port})$

$\text{score} = \text{uptime} / \text{distance}$




Subscriptions tree

The problem...


1. A node with a lot of subscribers **restarts its session**
2. How to **reestablish the network**?
 - Recursive node election?
 - Recursive random connections?
 - Preferential Attachment?

Maybe not! We must balance the total load on the central server and the connection reliability...

Random connected




Score based + Random




Tree **construction** and **broadcast** (2000 nodes)

Random tree

10 (ms)




16 rounds




Point-based tree + random attachment

16 (ms)



6 rounds

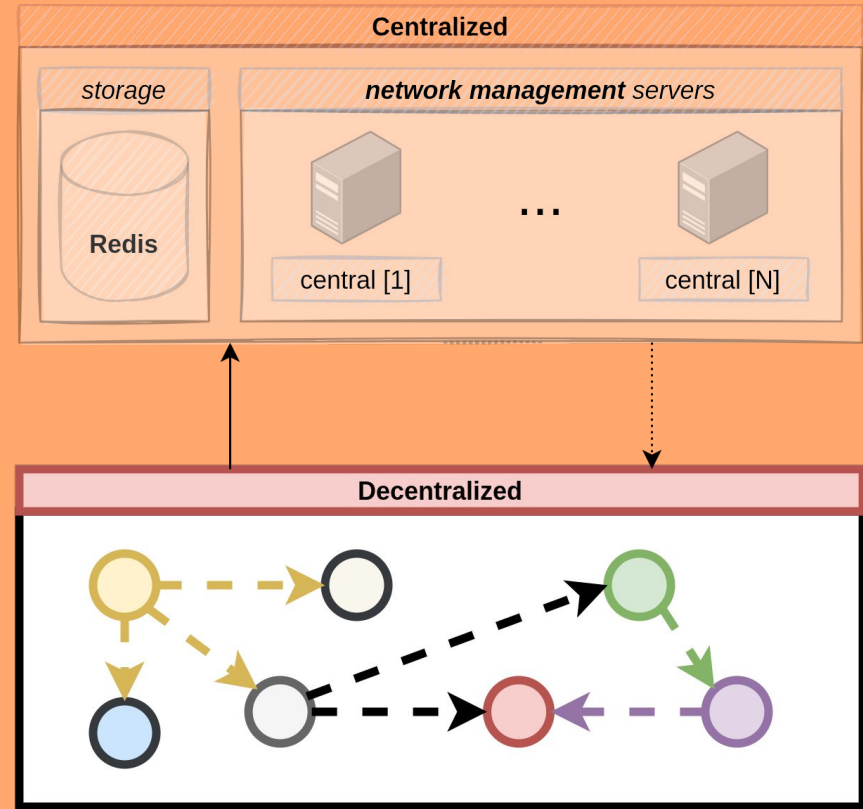


Decentralized portion

It's used for:

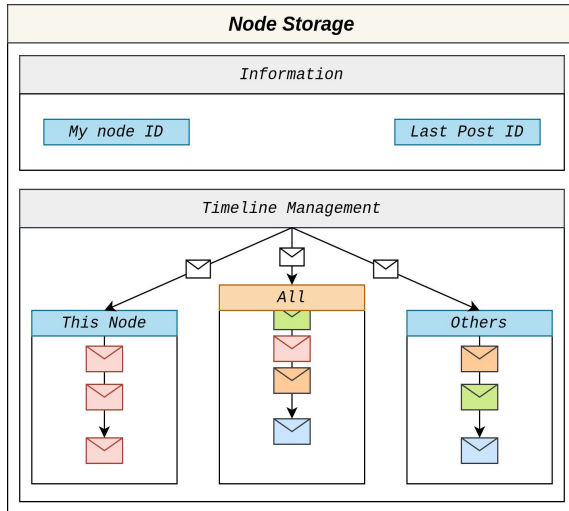
- a. Propagation of posts
- b. Timeline management

"Peer-to-Peer portion of the network"



Node information

Local storage



Use cases Services

1. Register

2. Login

a. gets the **elected candidates** ports for subscriptions and recovery

3. Logout

a. notifies central server that the node is **no longer online**

4. Subscription


a. node election

5. Timeline management

a. publish, forward, recover and post ordering

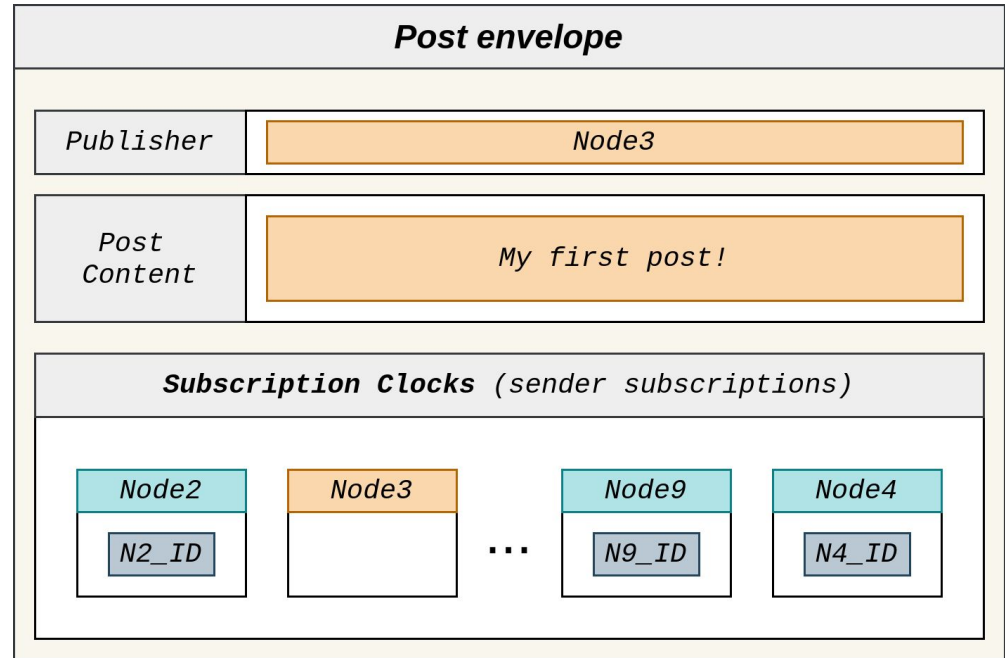
Timeline Management

- **Publish** a post
- **Reception** and **Forwarding**
- **Timeline Recovery** (for all subscriptions)
 - 1- Always after Login
 - 2- Recovers only from the last message of each subscription
 - 3- Some messages can be lost due to ephemerality



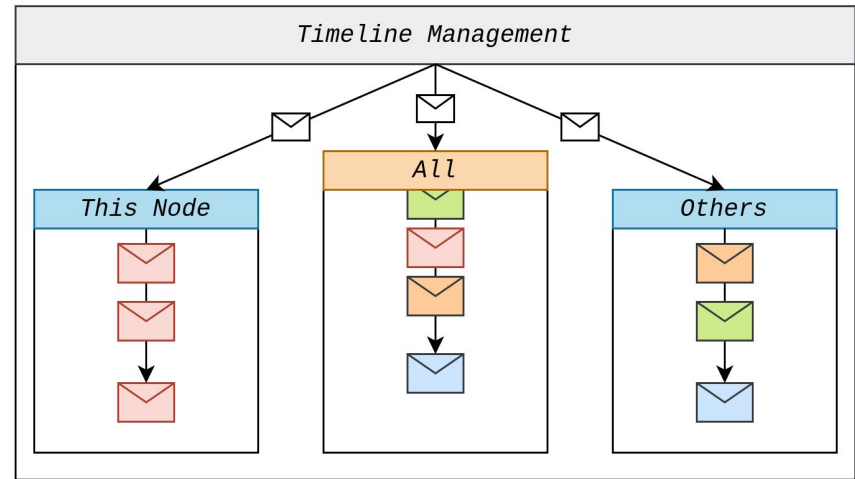
Post structure

- **Publisher** used to identify who sent the post
- **Post Content** that contains the message that the sender wants to deliver
- **Subscription Clocks** are important for ordering the messages



Message ordering

- We use the vector clock and order the posts **using the shared vector id's**
- Posts that are **concurrent** are distinguished using the **id of the sender**
- This solution can generate many paradoxical situations witch we didn't found a solution.



Architecture overview

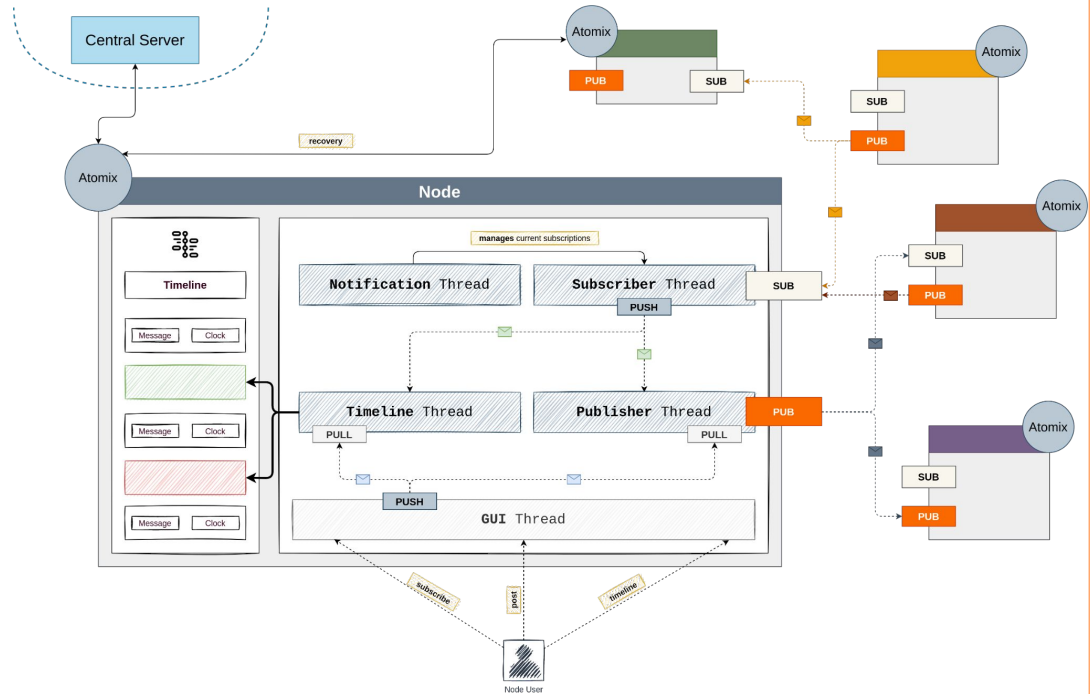
Java


ZeroMQ

Pub-Sub pattern
Push-Pull pattern

Atomix

Event-driven services





Conclusion

CAP Theorem

**Model
assumptions**



A84462 | Alexandre Miranda

A84961 | Alexandre Ferreira

A85227 | João Azevedo

A85315 | Miguel Cardoso

A85729 | Paulo Araújo